IndustriConnect: MCP Adapters and Mock-First Evaluation for AI-Assisted Industrial Operations

Anonymous Submission
Artifact repository: https://github.com/NordicAgents/IndustriConnect-MCPs

***Abstract*—AI assistants can decompose multi-step workflows, but they do not natively speak industrial protocols such as Modbus, MQTT/Sparkplug B, or OPC UA. This paper presents INDUSTRICONNECT, a suite of Model Context Protocol (MCP) adapters that expose industrial operations as schema-discoverable AI tools while preserving protocol-specific connectivity and safety controls. The system uses a common response envelope and a mock-first workflow so adapter behavior can be exercised locally before connecting to plant equipment. We evaluate the current flagship adapters for Modbus, MQTT/Sparkplug B, and OPC UA with deterministic mock environments and fixed tool workflows corresponding to 14 representative assistant tasks. Across 140 normal runs and 200 tool calls, all benchmarked tasks completed successfully; recovery experiments after forced endpoint restarts also recovered in all 30 trials. The results position INDUSTRICONNECT as a reproducible prototype for AI-assisted industrial operations, while also making clear that the evidence is limited to mock-based adapter behavior rather than production deployment.**

I. INTRODUCTION

Consider a simple operator request: check a temperature excursion on an OPC UA namespace, confirm a related Modbus register block, and publish a Sparkplug B status update for downstream systems. Each subtask is routine in isolation, but the overall workflow crosses incompatible operational technology (OT) protocols whose interfaces were not designed for modern AI assistants. This protocol gap remains a practical barrier to AI-assisted industrial automation even as Industry 4.0 programs push tighter coupling between information technology and plant-floor systems [1], [2].

Large language model (LLM) assistants work best with typed, discoverable tool interfaces rather than binary protocol frames or vendor SDKs. MCP provides a standard way for assistants to list tools, inspect schemas, and invoke operations through a common call interface [3]. MCP alone, however, does not solve OT connectivity, write-guarding, or the need to validate behavior safely before touching real equipment.

This paper reframes INDUSTRICONNECT as a bounded systems contribution: a prototype MCP-to-OT adapter architecture with a mock-first evaluation workflow. The paper does not claim a complete industrial platform, a new LLM reasoning method, or production deployment evidence. Instead, it focuses on the flagship adapters for Modbus, MQTT/Sparkplug B, and OPC UA, while retaining seven additional protocol modules as ecosystem breadth.

The manuscript makes three contributions:

1) A practical MCP-to-OT adapter architecture that standardizes tool discovery and response handling across heterogeneous industrial protocols.
2) A mock-first methodology for validating adapter behavior, write guards, and restart recovery without requiring physical devices.
3) A reproducible prototype evaluation on Modbus, MQTT/Sparkplug B, and OPC UA using deterministic workloads, latency instrumentation, and fault-injection trials.

II. DESIGN GOALS AND ARCHITECTURE

Heterogeneity. Industrial protocols expose different interaction models: register maps for Modbus, broker topics for MQTT, and browsable information models for OPC UA [4], [5]. The adapter layer must translate these differences into tool schemas that remain understandable to an AI client.

Schema discoverability. Each adapter registers MCP tools so a client can enumerate available operations and call them without bespoke protocol code. This keeps the assistant-facing interface uniform even when the underlying transport or addressing model differs.

Write safety. The adapters expose explicit write operations, but writes remain bounded by per-server configuration such as disable flags, address validation, and structured error returns. These controls matter because an AI-facing API must represent both successful writes and intended denials in a machine-readable way.

Mockability. Every flagship adapter is paired with a local mock so the full tool path can be exercised offline. This enables reproducible evaluation, debugging, and regression checks before any real controller or fieldbus is involved.

Figure 1 shows the core pattern. An MCP-compatible assistant issues tool calls to an INDUSTRICONNECT adapter. The adapter handles protocol translation, connection management, and any safety guards, then talks to either a mock endpoint or plant equipment. All benchmarked adapters return the same envelope, {`success`, `data`, `error`, `meta`}, where `meta` carries endpoint details and latency measurements. This does not erase protocol differences; instead, it gives a common wrapper around those differences so higher-level orchestration can remain consistent.

In the current prototype, this envelope is also the benchmark boundary. The adapters record wall-clock latency, retry attempts, and endpoint-specific metadata directly in the tool

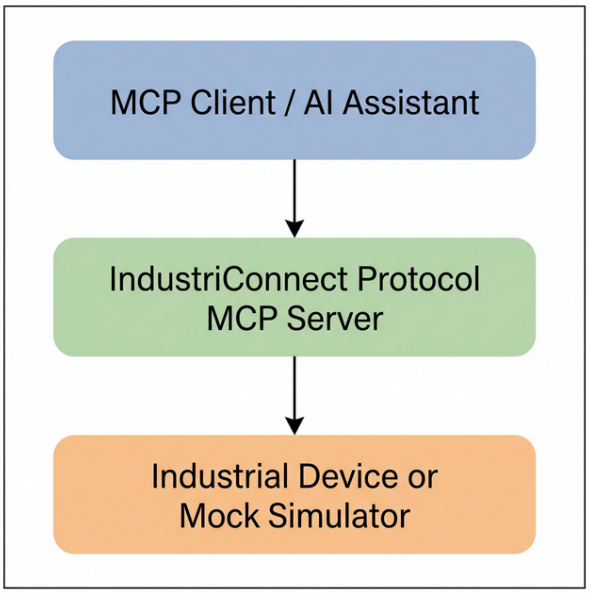

Fig. 1. INDUSTRICONNECT exposes industrial protocol clients as MCP servers so an AI assistant can invoke them through a common tool interface while targeting either mocks or real devices.

response. That makes cross-family comparison possible without pretending that a broker publish, a Modbus register read, and an OPC UA browse are identical operations.

III. PROTOTYPE IMPLEMENTATION

The repository currently contains ten protocol modules. Table I reports the tool inventory generated from registered handlers at benchmark time, not from manual counting. The three evaluated flagships are Modbus, MQTT/Sparkplug B, and OPC UA. The other seven modules demonstrate breadth but were not benchmarked in this paper and several remain roadmap-driven scaffolds around shared MCP patterns and local mocks. This distinction is deliberate: the paper claims measured behavior only for the flagship stacks.

Modbus. The Modbus adapter centers on register and coil operations, connection health checks, and guarded writes. In the evaluated workflow it handled register reads, write-and-readback verification, and explicit write denial when writes were disabled. The recovery benchmark also exercised reconnect behavior after a forced Modbus mock restart.

MQTT/Sparkplug B. The MQTT adapter exposes broker health, topic subscriptions, generic publishes, and Sparkplug B device-data publishes. Its MCP responses normalize broker connectivity and publish status into the common response envelope while preserving topic-level semantics. In practice this lets one assistant workflow mix ordinary MQTT control messages with Sparkplug-oriented telemetry updates.

OPC UA. The OPC UA adapter exposes node browsing, node reads, writes, and multi-node access against a local mock plant. Method invocation exists in the prototype but was excluded from the benchmark because the current mock method path is not yet stable enough for repeatable evaluation. The measured workload therefore focuses on repeatable assistant-facing operations instead of the full prototype surface.

TABLE I
PROTOCOL INVENTORY GENERATED FROM REGISTERED MCP TOOL HANDLERS.

Adapter	Tools	Maturity
Modbus	17	evaluated flagship
MQTT + Sparkplug B	15	evaluated flagship
OPC UA	7	evaluated flagship
BACnet/IP	7	roadmap/ scaffold
DNP3	8	roadmap/ scaffold
EtherCAT	14	roadmap/ scaffold
EtherNet/IP	17	roadmap/ scaffold
PROFIBUS DP/PA	11	roadmap/ scaffold
PROFINET	14	roadmap/ scaffold
Siemens S7comm	20	roadmap/ scaffold

TABLE II
BENCHMARK MATRIX USED IN THE MOCK-FIRST EVALUATION.

Family	Task set
Modbus	Ping; four-register read; register write/readback; guarded write denial
MQTT/SB	Broker info; subscribe to `sensors/#`; generic publish; Sparkplug B DDATA publish
OPC UA	Temperature read; actuator subtree browse; valve-position write/readback; full variable enumeration
Cross	Multi-adapter state snapshot; coordinated control sequence across Modbus, OPC UA, and MQTT

IV. EXPERIMENTAL SETUP

The evaluation targets adapter behavior rather than end-to-end LLM planning. To keep the workload reproducible, we fixed a set of representative assistant intents and executed them as deterministic MCP tool workflows against local mocks. This isolates connection handling, tool schemas, response normalization, and safety behavior from prompt variability [6].

Table II summarizes the workload. The normal benchmark matrix contains 14 tasks: four Modbus tasks, four MQTT/Sparkplug B tasks, four OPC UA tasks, and two cross-protocol tasks. The cross-protocol tasks exercise a multi-adapter state snapshot and a coordinated control sequence spanning all three flagship stacks. Each task was repeated 10 times, yielding 140 normal runs and 200 tool calls overall.

The workload includes one guarded-write scenario in which the Modbus adapter is configured read-only and is expected to return a structured denial rather than perform the write. Task success and tool-call success are therefore judged against a task-specific oracle, not simply against `success=true`. A structured write denial counts as correct behavior for that task.

The benchmark harness starts the local mocks, opens MCP stdio sessions for each adapter, executes the workload, and records JSON results from the shared response envelope. We

TABLE III
FAMILY-LEVEL BENCHMARK RESULTS. CROSS-PROTOCOL ROWS AGGREGATE THE TWO MULTI-ADAPTER TASKS ONLY.

Family	Tasks	Task (%)	Tool (%)	Med. (ms)	p95 (ms)	Rec. (s)
Modbus	4	100	100	1.47	3.62	4.64
MQTT/SB	4	100	100	1.76	2.30	3.45
OPC UA	4	100	100	6.05	114.86	6.12
Cross	2	100	100	6.58	7.73	–

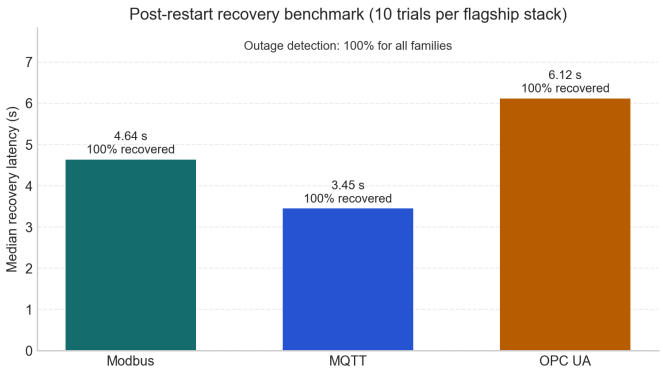

Fig. 2. Median end-to-end recovery latency after forced endpoint restarts. All three flagship stacks recovered in 10 of 10 trials and detected the outage in 10 of 10 trials.

measured task success, tool-call success, median latency, p95 latency, and retry counts from those records. We also injected one transient outage per flagship stack by stopping the mock or broker, confirming an operation failure, restarting the endpoint, and re-running a healthy operation. Each recovery scenario was repeated 10 times for a total of 30 restart trials.

V. RESULTS AND LIMITATIONS

Table III shows clean success rates across the full benchmark: 140 of 140 normal task executions succeeded, and all 200 tool calls matched their expected outcomes. No retries were required in the normal suite. Modbus and MQTT/Sparkplug B remained in the low single-digit millisecond range for both median and p95 latency. Cross-protocol workflows also completed reliably, showing that the common tool surface is usable for simple multi-adapter orchestration.

Task-level behavior is also informative. Modbus medians ranged from 1.19 ms for a simple ping to 2.85 ms for write-and-readback validation. MQTT/Sparkplug B stayed tightly grouped between 1.72 and 1.80 ms across broker inspection, subscription, and publish operations. Cross-protocol medians were 6.43 ms and 6.74 ms for the control and snapshot tasks, respectively.

OPC UA showed the largest tail latency. The baseline node read, browse, and write-readback tasks remained modest at 2.24 ms, 7.29 ms, and 4.50 ms median latency, but the variable-enumeration task reached 108.94 ms median and drove the family p95 to 114.86 ms. This is consistent with the information-model breadth of OPC UA rather than a generalized slowdown across all OPC UA operations.

Figure 2 summarizes the restart experiments. All three flagship stacks recovered in every trial after a forced outage. The measured recovery times are dominated by mock or broker restart time plus reconnection, not by the steady-state tool latency shown in Table III.

The failure modes also differed by stack. MQTT typically failed fast because the adapter could detect a disconnected broker immediately, while OPC UA recovery was slower because the failed read path waited on a longer timeout before the restarted server became reachable again. This suggests that future hardening work should tune timeout and reconnect policy per protocol rather than enforce a single global default.

Limitations. The evidence is intentionally narrow. This is a mock-only study, not a production deployment evaluation, not a cybersecurity validation, and not evidence that AI agents can safely perform closed-loop autonomous control on plant equipment. The benchmark also omits human baselines and excludes OPC UA method calls because the current prototype mock does not yet support them reliably enough for repeatable measurement.

VI. RELATED WORK

OT/IT convergence literature has long identified heterogeneous field protocols as a core integration problem for Industry 4.0 systems [1], [2]. OPC UA and MQTT are widely used normalization layers in modern industrial software, but each solves only part of the interoperability problem because brownfield plants still depend on multiple legacy and vendor-specific stacks [4], [5].

Recent LLM-agent work emphasizes reasoning plus external tool use rather than monolithic natural-language interfaces [6]. MCP contributes a practical, schema-driven interface for tool discovery and invocation [3]. Our contribution is not a new tool-use protocol; it is the adapter pattern, shared safety conventions, and benchmark evidence needed to map that protocol into industrial operations.

The mock-first workflow is adjacent to simulation and digital-twin practices used for industrial software validation [7]. In contrast to full digital-twin objectives, the focus here is narrower: verify adapter behavior, structured failures, and restart recovery before introducing real equipment.

VII. CONCLUSION

INDUSTRICONNECT is best understood as a prototype MCP-to-OT adapter architecture rather than a complete industrial platform. The current evidence shows that the Modbus, MQTT/Sparkplug B, and OPC UA adapters can execute deterministic mock workloads with consistent response semantics, guarded-write behavior, and repeatable restart recovery. The broader repository provides ten protocol modules, but only three are evaluated in this paper.

The next steps are straightforward: validate the same workflows on real hardware, harden the security and approval path around write-capable tools, and expand the benchmark beyond mock connectivity into operator-in-the-loop scenarios.

REFERENCES

[1] L. D. Xu, E. L. Xu, and L. Li, "Industry 4.0: state of the art and future trends," *International Journal of Production Research*, vol. 56, no. 8, pp. 2941–2962, 2018.

[2] M. Wollschlaeger, T. Sauter, and J. Jasperneite, "The future of industrial communication: Automation networks in the era of the internet of things and industry 4.0," *IEEE Industrial Electronics Magazine*, vol. 11, no. 1, pp. 17–27, 2017.

[3] Anthropic, "Model context protocol specification," https://modelcontextprotocol.io, 2024, accessed 2026-03-12.

[4] S. Cavalieri and F. Chiacchio, "Analysis of OPC UA performances," *Computer Standards & Interfaces*, vol. 36, no. 1, pp. 165–177, 2013.

[5] B. Mishra and A. Kertesz, "The use of MQTT in M2M and IoT systems: A survey," *IEEE Access*, vol. 8, pp. 201 071–201 086, 2020.

[6] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao, "ReAct: Synergizing reasoning and acting in language models," *arXiv preprint arXiv:2210.03629*, 2022. [Online]. Available: https://arxiv.org/abs/2210.03629

[7] A. Fuller, Z. Fan, C. Day, and C. Barlow, "Digital twin: Enabling technologies, challenges and open research," *IEEE Access*, vol. 8, pp. 108 952–108 971, 2020.