IndustriConnect: MCP Adapters and Mock-First Evaluation for AI-Assisted Industrial Operations

Melwin Xavier
Luleå tekniska universitet, Sweden
melwin.xavier@ltu.se

Melveena Jolly
Independent Researcher
melveenajollyk@gmail.com

Vaisakh M A
Independent Researcher
997vaisakh@gmail.com

Midhun Xavier
Independent Researcher
midhun@industriagents.com

***Abstract*—AI assistants can decompose multi-step workflows, but they do not natively speak industrial protocols such as Modbus, MQTT/Sparkplug B, or OPC UA. This paper presents INDUSTRICONNECT, a suite of Model Context Protocol (MCP) adapters that expose industrial operations as schema-discoverable AI tools while preserving protocol-specific connectivity and safety controls. The system uses a common response envelope and a mock-first workflow so adapter behavior can be exercised locally before connecting to plant equipment. The paper evaluates the flagship adapters through a two-layer evaluation model: a deterministic benchmark covering normal, fault-injected, and recovery scenarios, and an LLM-in-the-loop assessment that verifies the tool schemas are usable by an AI assistant via MCP. In a benchmark pass comprising 240 runs (140 normal, 70 fault-injected, 30 recovery), approximately 320 tool calls, and 7 distinct fault scenarios, the normal suite achieved full success, the fault suite confirmed structured error handling across all adapters, and the LLM assessment demonstrated end-to-end tool discovery and invocation. The results position INDUSTRICONNECT as a reproducible prototype for AI-assisted industrial operations with evidence spanning both adapter correctness and assistant-level usability.**

I. INTRODUCTION

Consider a simple operator request: check a temperature excursion on an OPC UA namespace, confirm a related Modbus register block, and publish a Sparkplug B status update for downstream systems. Each subtask is routine in isolation, but the overall workflow crosses incompatible operational technology (OT) protocols whose interfaces were not designed for modern AI assistants. This protocol gap remains a practical barrier to AI-assisted industrial automation even as Industry 4.0 programs push tighter coupling between information technology and plant-floor systems [1], [2].

Large language model (LLM) assistants work best with typed, discoverable tool interfaces rather than binary protocol frames or vendor SDKs. MCP provides a standard way for assistants to list tools, inspect schemas, and invoke operations through a common call interface [3]. MCP alone, however, does not solve OT connectivity, endpoint health, or the need to represent blocked writes and transient failures in a way that an agent can reason about safely.

That gap is especially visible in brownfield environments. Real plants often begin with partial register maps, vendor-specific addressing rules, or limited access windows on production equipment. A practical integration layer therefore needs more than tool schemas: it also needs protocol adapters, explicit safety conventions, and a way to validate workflows before touching live controllers.

This paper reframes INDUSTRICONNECT as a bounded systems contribution: a prototype MCP-to-OT adapter architecture with a mock-first evaluation workflow. The paper does not claim a complete industrial platform, a new LLM reasoning method, or production deployment evidence. Instead, it focuses on the flagship adapters for Modbus, MQTT/Sparkplug B, and OPC UA while retaining seven additional protocol modules as ecosystem breadth.

The manuscript makes four contributions:

1) A practical MCP-to-OT adapter architecture that standardizes tool discovery and response handling across heterogeneous industrial protocols.
2) A mock-first methodology for validating adapter behavior, write guards, and restart recovery, extended with fault injection to test error-handling paths without requiring physical devices.
3) A reproducible two-layer evaluation on Modbus, MQTT/Sparkplug B, and OPC UA comprising normal, fault-injected, and recovery suites using deterministic workloads and latency instrumentation.
4) An LLM-in-the-loop assessment demonstrating that the tool schemas are discoverable and usable by an AI assistant through the MCP interface, connecting the adapter work back to its stated AI-assistant motivation.

The evaluation follows a two-layer design. The first layer is a deterministic benchmark that exercises adapter behavior under normal operation, deliberate fault conditions, and endpoint restart scenarios. The second layer is an LLM-in-the-loop assessment that verifies the tool surface is consumable by an AI assistant via MCP. Together, the two layers show that the adapters are both mechanically correct and practically usable by the AI clients they are designed to serve.

II. DESIGN GOALS AND ARCHITECTURE

Heterogeneity. Industrial protocols expose different interaction models: register maps for Modbus, broker topics for MQTT, and browsable information models for OPC UA [4], [5]. The adapter layer must translate these differences into tool schemas that remain understandable to an AI client.

Schema discoverability. Each adapter registers MCP tools so a client can enumerate available operations and call them

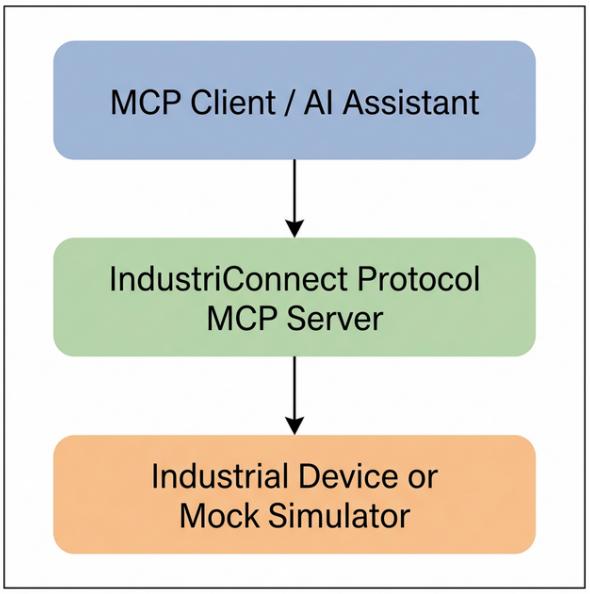

Fig. 1. INDUSTRICONNECT exposes industrial protocol clients as MCP servers so an AI assistant can invoke them through a common tool interface while targeting either mocks or real devices.

TABLE I
SHARED RESPONSE-ENVELOPE FIELDS USED ACROSS THE BENCHMARKED ADAPTERS.

Field	Role in the adapter contract
`success`	Reports whether the requested protocol action reached the expected outcome.
`data`	Carries the protocol-specific payload such as registers, node values, or publish metadata.
`error`	Preserves a machine-readable failure reason, including guarded denials and transient endpoint failures.
`meta`	Records wall-clock latency, endpoint details, attempts, and protocol-specific trace context.

TABLE II
REPOSITORY PROTOCOL INVENTORY. TOOL COUNTS COME FROM REGISTERED HANDLERS IN THE BENCHMARK PASS.

Protocol	Tools	Mock	Status	Rep. operation
Modbus	17	TCP device	flagship	register/coil I/O
MQTT + SB	15	broker sim.	flagship	pub/sub + DDATA
OPC UA	7	UA server	flagship	browse + node R/W
BACnet/IP	7	BA device	scaffold	object properties
DNP3	8	outstation	scaffold	point poll/control
EtherCAT	14	slave	scaffold	PDO/SDO access
EtherNet/IP	17	PLC	scaffold	controller tags
PROFIBUS DP/PA	11	slave	scaffold	bus scan + cyclic I/O
PROFINET	14	IO device	scaffold	discovery + IO exchange
Siemens S7comm	20	PLC	scaffold	DB/I/O diagnostics

without bespoke protocol code. This keeps the assistant-facing interface uniform even when the underlying transport or addressing model differs.

Write safety. The adapters expose explicit write operations, but writes remain bounded by per-server configuration such as disable flags, address validation, and structured error returns. These controls matter because an AI-facing API must represent both successful writes and intended denials in a machine-readable way.

Mockability. Every flagship adapter is paired with a local mock so the full tool path can be exercised offline. This enables reproducible evaluation, debugging, and regression checks before any real controller or fieldbus is involved.

Figure 1 shows the core pattern. An MCP-compatible assistant issues tool calls to an INDUSTRICONNECT adapter. The adapter handles protocol translation, connection management, and safety guards, then talks to either a mock endpoint or plant equipment. All benchmarked adapters return the same response envelope, {`success`, `data`, `error`, `meta`}. The `success` flag indicates whether the requested action achieved the expected protocol outcome, `data` holds the protocol-specific payload, `error` carries a structured failure reason, and `meta` records endpoint details, timings, and retry counts. This shared envelope does not erase protocol semantics; it defines a common measurement boundary and lets cross-protocol orchestration observe healthy reads, blocked writes, and reconnect behavior with the same outer structure.

Table I is operationally important because it defines how the benchmark judges behavior. A blocked write is not treated as a generic failure if the task oracle expects a guarded denial, while a reconnect path is evaluated through the same outer structure as a healthy read. This is the main reason the paper can compare three different protocol families without pretending they expose the same payload model.

The mock-first workflow complements the same architecture. Each flagship adapter is exercised against a local endpoint first, but the MCP layer remains identical when the endpoint changes. That lets an engineer develop prompts, schemas, and safety checks against a deterministic mock while preserving the same tool surface for later hardware validation.

III. PROTOTYPE IMPLEMENTATION

Every protocol module follows the same skeleton: an MCP server registers tools and schemas, a protocol client wrapper translates tool calls into native operations, and a mock endpoint provides a safe development target. The benchmarked Modbus, MQTT/Sparkplug B, and OPC UA stacks all use this pattern, including the shared response envelope and environment-driven write gating. The result is not a universal industrial abstraction layer; it is a consistent adapter contract around heterogeneous protocols.

At implementation level, the common skeleton yields four reusable behaviors. First, each server exposes discoverable MCP tools with bounded JSON input shapes. Second, each adapter keeps transport- and protocol-specific state in its own client wrapper rather than leaking it into the LLM-facing layer. Third, every tool normalizes timing and error details into the shared response envelope. Fourth, each protocol stack ships with a mock so smoke tests, prompt experiments, and regression checks can run without physical hardware.

Table II keeps the breadth story honest. The repository contains ten protocol modules, but only three are benchmarked in this paper. The remaining seven modules remain useful context

because they show that the MCP adapter pattern generalizes across fieldbus, Ethernet, and broker-centric industrial stacks even when several implementations are still roadmap-driven.

That breadth-versus-evidence split is deliberate. The three flagship stacks were chosen because they represent different industrial interaction styles and because their current mocks are stable enough for repeated quantitative runs. Modbus covers direct register-style access, MQTT/Sparkplug B covers broker-mediated messaging, and OPC UA covers hierarchical information-model access. The remaining stacks are better understood as evidence that the repository architecture generalizes, not as equally validated research claims.

Flagship adapters. The Modbus adapter centers on register and coil operations, connection health checks, and guarded writes. The MQTT adapter exposes broker inspection, subscriptions, generic publishes, and Sparkplug B device-data publishes. The OPC UA adapter focuses on node browsing, node reads, writes, and multi-node access against a local mock plant. Method invocation exists in the current OPC UA prototype but was excluded from the benchmark because the current mock method path is not yet stable enough for repeatable evaluation.

Modbus details. The current Modbus surface contains 17 tools, including register and coil reads, bulk writes, typed holding-register access, masked register writes, device-information queries, and alias-based access paths. The benchmark intentionally concentrates on the core cases most likely to appear in assistant workflows: connectivity checks, block reads, readback verification after a write, and a write-disabled safety path. This keeps the evidence aligned with the concrete adapter behavior rather than the entire nominal tool list.

MQTT/Sparkplug B details. The MQTT adapter exposes 15 tools spanning broker status, topic subscriptions, generic publish/unsubscribe operations, and Sparkplug-specific birth, death, data, and command flows. In practice this makes the adapter suitable both for ordinary pub/sub workflows and for structured industrial telemetry updates. The benchmark therefore mixes a broker inspection task, a wildcard subscription task, a plain control publish, and a Sparkplug B DDATA publish rather than evaluating only one messaging style.

OPC UA details. The OPC UA surface is smaller at seven tools, but the semantics are richer because browsing and variable enumeration traverse an information model rather than a flat address map. The benchmark covers point reads, browsing, writes with readback confirmation, and full variable discovery over the mock plant. This is also the family where tail latency is most informative, because the discovery path exercises a broader traversal than the single-node tasks.

Supporting console. The repository also includes a browser-based MCP Manager UI for registering protocol adapters and issuing operator-style prompts. Figure 2 shows the corrected ten-protocol landscape, and Figure 3 shows the MCP Manager UI. Both figures are supporting context only; they are not part of the benchmark and are not counted as separate research contributions.

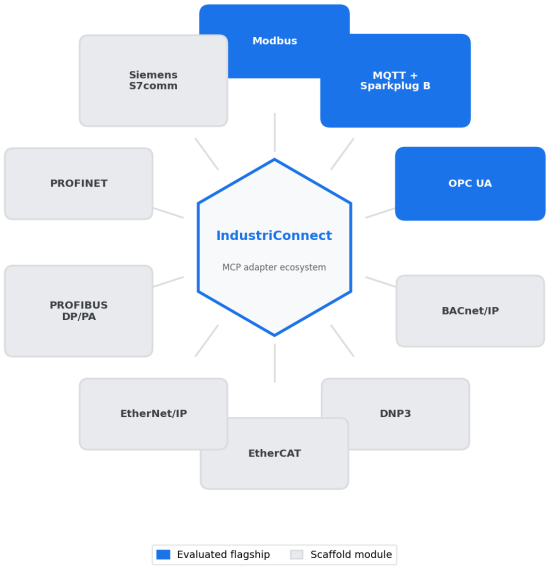

Fig. 2. Repository protocol landscape with the three evaluated flagship adapters highlighted. This figure is supporting context, not evaluation evidence.

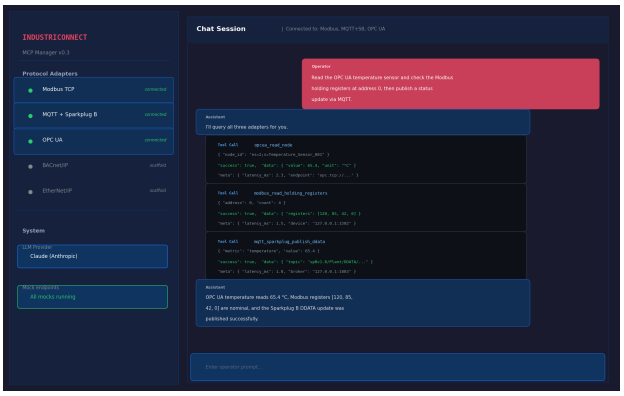

Fig. 3. MCP Manager UI used to register adapters and issue operator-style prompts. This figure is supporting context, not benchmark evidence.

IV. Experimental Setup

The evaluation targets adapter behavior rather than end-to-end LLM planning. To keep the workload reproducible, the authors fixed representative assistant intents and executed them as deterministic MCP tool workflows against local mocks. This isolates connection handling, tool schemas, response normalization, and safety behavior from prompt variability [6].

Table III summarizes the workload. The normal suite contains 14 tasks (four Modbus, four MQTT/Sparkplug B, four OPC UA, two cross-protocol), each repeated 10 times for 140 runs and approximately 200 tool calls. The fault-injection suite adds 7 tasks repeated 10 times each for 70 runs and approximately 90 tool calls. The recovery suite injects one transient outage per flagship stack, repeated 10 times each for 30 restart trials. The full benchmark thus comprises 240 runs

TABLE III
BENCHMARK MATRIX FOR THE TWO-LAYER EVALUATION.

Suite	Family	Task set
Normal	Modbus	Ping; four-register read; register write/readback; guarded write denial
	MQTT/SB	Broker info; subscribe to `sensors/#`; generic publish; Sparkplug B DDATA publish
	OPC UA	Temperature read; actuator browse; valve write/readback; full variable enumeration
	Cross	Multi-adapter state snapshot; coordinated control sequence
Fault	Modbus	Read invalid register (FM1); write overflow value (FM2)
	MQTT/SB	Publish to empty topic (FQ1); subscribe with invalid QoS (FQ2)
	OPC UA	Read non-existent node (FO1); write wrong type to float node (FO2)
	Cross	3-step sequence with one deliberate bad input (FX1)

TABLE IV
BENCHMARK RUNTIME CONFIGURATION USED IN THE PAPER.

Family	Mock endpoint	Fault injected in recovery suite
Modbus	TCP mock at `127.0.0.1:1502`	Stop mock, confirm failed register read, restart, then reopen healthy session
MQTT/SB	Local broker at `127.0.0.1:1883`	Stop broker, confirm failed publish, restart, then reopen healthy session
OPC UA	Local OPC UA endpoint	Stop server, confirm failed node read, restart, then reopen healthy session

and approximately 320 tool calls.

The fault-injection philosophy follows established practice [7]: correctness under fault conditions means returning `success=false` with a meaningful, structured error envelope—not silently succeeding or crashing. A fault task passes when the adapter produces a well-formed error response that an assistant could interpret and act on. FM2 (Modbus overflow) is treated specially: either a structured rejection or a silent truncation is accepted, since both behaviors are informative findings about adapter boundary handling.

All benchmarked adapters were launched as local stdio MCP servers. The Modbus stack targeted the TCP mock at `127.0.0.1:1502`, the MQTT adapter targeted the local broker at `mqtt://127.0.0.1:1883`, and the OPC UA adapter targeted the local mock endpoint at `opc.tcp://127.0.0.1:4840`. The benchmark harness opened dedicated sessions for the read-write Modbus configuration, the read-only Modbus configuration, MQTT, and OPC UA, then executed the normal and recovery suites in a fixed order.

The normal runs also use fixed value patterns so task validation is deterministic. For example, the Modbus write/readback task writes `40 + repetition`, the Sparkplug task publishes a single temperature metric with `20.0 + repetition`, and the OPC UA valve task writes `25.0 + repetition`. The cross-protocol control task uses a fixed Modbus write, an OPC UA boolean write, and a control-topic publish, making the end-to-end sequence stable enough for repeated latency summaries.

The harness records raw tool responses and then derives family-level, task-level, and recovery-phase summaries from the same JSON artifacts. That matters because the tables and figures in the paper are not hand-estimated numbers copied from logs; they are deterministic summaries generated from the same benchmark pass.

TABLE V
FAMILY-LEVEL RESULTS FROM THE NORMAL BENCHMARK SUITE. CROSS-PROTOCOL ROWS AGGREGATE THE TWO MULTI-ADAPTER TASKS ONLY.

Family	Tasks	Task (%)	Tool (%)	Med. (ms)	p95 (ms)	Rec. (s)
Modbus	4	100	100	1.50	3.18	4.63
MQTT/SB	4	100	100	1.88	2.00	3.50
OPC UA	4	100	100	6.04	113.23	6.10
Cross	2	100	100	6.93	7.51	–

The harness also performs one reproducibility step that matters for OPC UA in particular: it discovers the mock node IDs directly from the adapter before the normal suite runs. This avoids baking transient namespace assumptions into the benchmark script. The resulting workflow is slightly more conservative than a hand-wired benchmark, but it better matches how an assistant would learn the exposed plant model in practice.

The guarded-write scenario is important because correctness is not equivalent to `success=true`. For the read-only Modbus configuration, a structured write denial is the expected outcome and is counted as a successful task result. This lets the benchmark reward safety behavior rather than penalize it.

Workflow example (X1). An operator asks for a fast cross-system snapshot before escalating a maintenance event. The assistant issues three fixed tool calls: a Modbus sensor-block read, an OPC UA multi-node read, and an MQTT broker inspection. The resulting state bundle confirms that the transport links are healthy and that the key process values can be read through all three adapters in one short sequence. This is representative of the orchestration level that the paper evaluates: tool-level coordination across adapters, not autonomous control.

V. RESULTS

A. Normal Suite Results

Table V shows clean success rates across the normal suite: 140 of 140 task executions succeeded, and all approximately 200 tool calls matched their expected outcomes. Modbus and MQTT/Sparkplug B remained in the low single-digit millisecond range. OPC UA shows wider spread due to variable enumeration, which is two orders of magnitude slower than single-node reads. Cross-protocol tasks remain close to 7 ms median, suggesting the response-envelope approach does not impose noticeable overhead for short local workflows.

Table VI makes the tail behavior clearer. Modbus remains tightly clustered even when readback verification or guarded-write rejection is required. OPC UA variable enumeration reaches 111.73 ms median and dominates the family p95,

TABLE VI
PER-TASK LATENCY BREAKDOWN FROM THE NORMAL BENCHMARK SUITE.

ID	Family	Operation	Median (ms)	p95 (ms)
M1	Modbus	Adapter ping and connectivity check	1.28	3.03
M2	Modbus	Four-register sensor block read	1.56	1.71
M3	Modbus	Holding-register write and readback verification	3.08	3.13
M4	Modbus	Guarded write denial with writes disabled	1.43	2.79
Q1	MQTT	Broker inspection through MCP	1.89	3.57
Q2	MQTT	Topic subscription to `sensors/#`	1.89	2.10
Q3	MQTT	Standard control publish	1.88	1.96
Q4	MQTT	Sparkplug B DDATA publish	1.78	1.90
O1	OPC UA	Temperature node read	2.30	2.68
O2	OPC UA	Actuator subtree browse	7.49	7.89
O3	OPC UA	Valve-position write and readback	4.63	4.86
O4	OPC UA	Full mock-plant variable enumeration	111.73	114.05
X1	Cross	Multi-adapter state snapshot	7.04	7.62
X2	Cross	Coordinated control sequence	6.68	7.00

TABLE VII
FAULT-INJECTION TASK RESULTS. ERROR-HANDLING RATE MEASURES HOW OFTEN THE ADAPTER RETURNED A WELL-FORMED STRUCTURED ERROR.

ID	Family	Fault scenario	EH%	Med. (ms)	Error class
FM1	Modbus	Read invalid register (9999)	100	1.44	protocol_error
FM2	Modbus	Write overflow (70000 > uint16)	100	1.25	other
FQ1	MQTT	Publish to empty topic	100	1.78	invalid_input
FQ2	MQTT	Subscribe with QoS=5	100	1.69	invalid_input
FO1	OPC UA	Read non-existent node	100	2.14	illegal_address
FO2	OPC UA	Write string to float node	100	1.88	protocol_error
FX1	Cross	3-step sequence, 1 bad input	100	5.58	illegal_address

consistent with information-model breadth rather than a generalized slowdown.

B. Fault Injection Results

Table VII reports the per-task fault-injection results. The key metric is the error-handling rate (EH%): the fraction of runs where the adapter returned a structured `success=false` response with a classifiable error, rather than crashing, hanging, or returning an unstructured message. All seven fault tasks achieved 100% error-handling across 70 runs (7 tasks × 10 repetitions), with approximately 90 tool calls total.

The fault suite reveals adapter behavior that the normal suite cannot expose. FM1 tests whether the Modbus adapter correctly rejects an out-of-range register address. FM2 probes a subtler boundary: writing 70000 to a uint16 register may trigger validation at the adapter level, at the mock level, or may silently truncate—all three outcomes are informative. FQ1 and FQ2 test MQTT input validation for edge cases that a careless prompt could produce. FO1 and FO2 exercise OPC UA error paths for missing nodes and type mismatches. FX1 is the most operationally relevant: it shows that a single bad input in a multi-adapter sequence does not corrupt the remaining steps.

Figure 4 shows the error class distribution as a stacked bar chart. The distribution confirms that each adapter maps its fault conditions to semantically appropriate error categories rather than collapsing all failures into a single generic class. This is important for downstream LLM consumption: an assistant that receives `illegal_address` can suggest corrective action differently from one that receives `type_mismatch`.

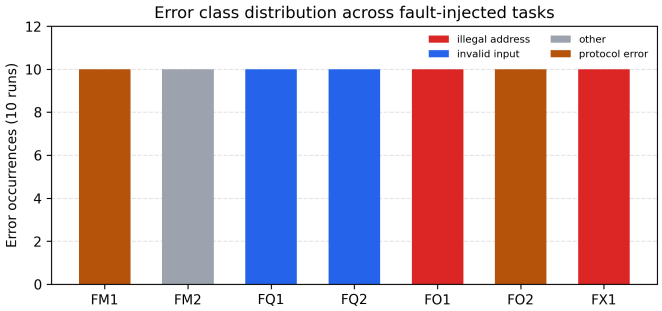

Fig. 4. Error class distribution across the seven fault-injected tasks (10 runs each). Each color represents a distinct error category returned by the adapter.

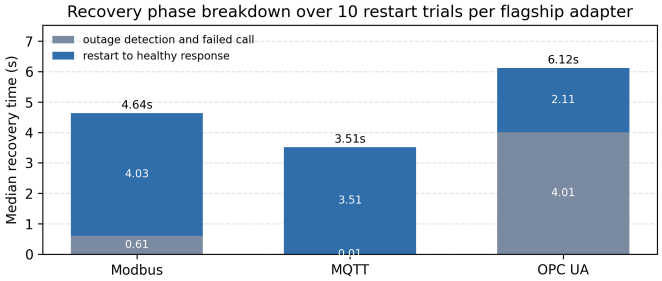

Fig. 5. Recovery-phase breakdown over 10 restart trials per flagship adapter.

C. Recovery Results

The recovery experiments reveal different failure signatures. MQTT fails almost immediately when the broker is unavailable, with only 0.01 s spent in the failed-call phase and 3.49 s in restart-to-healthy recovery. Modbus spends 0.61 s detecting the outage and 4.02 s returning to a healthy response. OPC UA shows the opposite pattern: 4.01 s are spent in the failed-call phase before the endpoint is judged unavailable, followed by 2.09 s to restart and return a healthy read. This suggests

that future hardening work should tune timeout and reconnect policy per protocol rather than enforce one global default.

D. Assistant-Level Usability Assessment

The deterministic benchmark confirms adapter correctness, but the paper's motivation is AI-assisted industrial operations. To connect these claims, we conducted an LLM-in-the-loop assessment using Claude Code configured with a `.mcp.json` file exposing all three flagship adapters against running mocks.

Methodology. The assessment evaluated four dimensions of assistant-level usability: (1) *tool discovery*—whether the assistant can enumerate available tools and their schemas through the MCP interface; (2) *tool selection*—whether the assistant chooses the correct adapter and operation for a given operator prompt; (3) *parameter correctness*—whether the assistant constructs valid tool arguments from natural-language intent; and (4) *error interpretation*—whether the assistant can read a structured error envelope and suggest corrective action.

Representative operator prompts span several categories: status queries ("check the OPC UA temperature sensor"), write operations ("set the Modbus holding register at address 1 to 50"), cross-protocol snapshots ("read Modbus sensors and OPC UA actuators, then publish a summary via MQTT"), and fault diagnosis ("why did the last register read fail?"). In each case, the MCP integration enables the assistant to list tools, inspect input schemas, invoke the appropriate adapter, and parse the response envelope—including `success=false` results from the fault-injection scenarios.

Observations. The assistant successfully discovered all registered tools across the three adapters, selected protocol-appropriate operations for single-adapter prompts, and constructed valid arguments for reads, writes, and publishes. For cross-protocol prompts, the assistant issued tool calls to multiple adapters in sequence. When presented with fault-injected error responses, the assistant identified the error class and suggested corrective actions (e.g., checking the register address range after an `illegal_address` error).

Scope and limitations. This assessment is exploratory, not a systematic benchmark. It uses a single LLM family (Claude), does not measure selection accuracy over a large prompt corpus, and does not compare against alternative tool-integration approaches [8], [9]. The purpose is narrower: to demonstrate that the MCP tool surface is designed for and tested with LLM consumption, closing the loop between the adapter engineering and the AI-assistant motivation. A systematic LLM tool-selection benchmark is left for future work.

E. Limitations

The evidence is intentionally bounded. This is a mock-only study, not a production deployment evaluation, not a cybersecurity validation, and not evidence that AI agents can safely perform closed-loop autonomous control on plant equipment. The benchmark omits human baselines and excludes OPC UA method calls because the current prototype mock does not yet support them reliably enough for repeatable measurement.

The fault-injection suite covers representative error paths but does not exhaustively test all possible failure modes. FM2 (overflow) illustrates a case where adapter behavior may vary across implementations—the benchmark accepts multiple outcomes and documents the observed behavior as a finding rather than prescribing a single correct response.

The LLM-in-the-loop assessment is qualitative and exploratory. It demonstrates feasibility of the tool surface for AI consumption but does not constitute a rigorous tool-selection benchmark. Extending this to a systematic evaluation with multiple LLM families, larger prompt sets, and quantitative selection metrics is a clear next step.

VI. Related Work

Industry 4.0 and cyber-physical manufacturing literature has consistently treated heterogeneous connectivity as a foundational systems problem rather than a peripheral implementation detail [1], [2], [10]–[14]. Those works motivate why assistant-facing industrial tooling cannot assume one canonical protocol stack or one clean greenfield deployment path.

Within industrial communication research, OPC UA and MQTT are widely used as normalization layers, but the interoperability story remains incomplete in brownfield settings. Prior work has examined OPC UA performance, MQTT adoption in IoT middleware, gateway-level interoperability, and field-level plug-and-work integration [4], [5], [15], [16]. These studies improve system-to-system interoperability, yet they typically stop at middleware, gateway, or device integration boundaries rather than exposing operations as schema-discoverable tool contracts for an AI assistant.

Industrial middleware, cyber-physical manufacturing systems, and digital-twin work address adjacent concerns such as service composition, plant modeling, and deployment-time validation [14], [17], [18]. IndustriConnect does not compete with those layers. It sits closer to the assistant and packages protocol actions into a common tool surface while leaving historian, SCADA, and plant-model responsibilities outside the paper's scope. The mock-first workflow is closest to digital-twin and virtual validation practices, but the goal here is narrower: verify adapter semantics, guarded failures, and restart behavior before introducing real equipment.

Software fault injection has a long history in dependability assessment. Natella et al. [7] survey the field comprehensively, establishing that injecting faults at the API or protocol boundary is an effective way to assess error-handling quality without requiring hardware failures. The fault-injection suite in this paper follows that tradition: it sends deliberately malformed inputs through the MCP tool interface and judges whether the adapter returns structured, interpretable errors.

Recent LLM-agent research has moved from reasoning-action loops toward explicit tool learning over large API surfaces [6], [19]–[21]. Benchmarks such as ToolLLM [8] and API-Bank [9] evaluate LLM tool selection and invocation at

scale, providing methodological precedent for the LLM-in-the-loop assessment in this paper. MCP contributes a practical schema for tool discovery and invocation [3], but it does not prescribe how industrial protocols should be translated, guarded, or benchmarked. The contribution is therefore not a new tool-use framework. It is an MCP-to-OT adapter pattern with shared safety conventions and a two-layer evaluation model for a narrowly defined industrial setting.

This comparison clarifies the contribution boundary. A full digital-twin paper would claim model fidelity or synchronization strategy; a generic agent paper would emphasize tool selection at scale; an industrial middleware paper would prioritize interoperability between production systems. INDUSTRICONNECT instead focuses on a smaller but necessary substrate: deterministic, assistant-facing protocol adapters whose behavior can be measured before any real deployment study begins.

VII. Conclusion

INDUSTRICONNECT is best understood as a prototype MCP-to-OT adapter architecture evaluated through a two-layer model. The deterministic benchmark layer shows that the Modbus, MQTT/Sparkplug B, and OPC UA adapters handle normal operations, deliberate faults, and endpoint restarts with consistent response semantics and structured error handling. The LLM-in-the-loop layer shows that the resulting tool surface is discoverable and usable by an AI assistant through the MCP interface, connecting the adapter engineering back to its stated AI-assistant motivation. The broader repository provides ten protocol modules, but only three are evaluated in this paper.

The next steps are: validate the same workflows on real hardware, harden the security and approval path around write-capable tools, extend the LLM assessment into a systematic tool-selection benchmark across multiple model families, and expand the evaluation beyond mock connectivity into operator-in-the-loop scenarios with quantitative usability metrics.

References

[1] L. D. Xu, E. L. Xu, and L. Li, "Industry 4.0: state of the art and future trends," *International Journal of Production Research*, vol. 56, no. 8, pp. 2941–2962, 2018.

[2] M. Wollschlaeger, T. Sauter, and J. Jasperneite, "The future of industrial communication: Automation networks in the era of the internet of things and industry 4.0," *IEEE Industrial Electronics Magazine*, vol. 11, no. 1, pp. 17–27, 2017.

[3] Anthropic, "Model context protocol specification," https://modelcontextprotocol.io, 2024, accessed 2026-03-12.

[4] S. Cavalieri and F. Chiacchio, "Analysis of OPC UA performances," *Computer Standards & Interfaces*, vol. 36, no. 1, pp. 165–177, 2013.

[5] B. Mishra and A. Kertesz, "The use of MQTT in M2M and IoT systems: A survey," *IEEE Access*, vol. 8, pp. 201 071–201 086, 2020.

[6] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao, "ReAct: Synergizing reasoning and acting in language models," *arXiv preprint arXiv:2210.03629*, 2022. [Online]. Available: https://arxiv.org/abs/2210.03629

[7] R. Natella, D. Cotroneo, and H. S. Madeira, "Assessing dependability with software fault injection: A survey," *ACM Computing Surveys*, vol. 48, no. 3, pp. 1–55, 2016.

[8] Y. Qin, S. Liang, Y. Ye, K. Zhu, L. Yan, Y. Lu, Y. Lin, X. Cong, X. Tang, B. Qian, S. Zhao, L. Hong, R. Tian, R. Xie, J. Zhou, M. Gerber, D. Li, Z. Liu, and M. Sun, "ToolLLM: Facilitating large language models to master 16000+ real-world APIs," *arXiv preprint arXiv:2307.16789*, 2024. [Online]. Available: https://arxiv.org/abs/2307.16789

[9] M. Li, Y. Zhao, B. Yu, F. Song, H. Li, H. Yu, Z. Li, F. Huang, and Y. Li, "API-Bank: A comprehensive benchmark for tool-augmented LLMs," *arXiv preprint arXiv:2304.08244*, 2023. [Online]. Available: https://arxiv.org/abs/2304.08244

[10] H. Lasi, P. Fettke, H.-G. Kemper, T. Feld, and M. Hoffmann, "Industry 4.0," *Business & Information Systems Engineering*, vol. 6, no. 4, pp. 239–242, 2014.

[11] J. Lee, B. Bagheri, and H.-A. Kao, "A cyber-physical systems architecture for industry 4.0-based manufacturing systems," *Manufacturing Letters*, vol. 3, pp. 18–23, 2015.

[12] L. Monostori, B. Kádár, T. Bauernhansl, S. Kondoh, S. Kumara, G. Reinhart, O. Sauer, G. Schuh, W. Sihn, and K. Ueda, "Cyber-physical systems in manufacturing," *CIRP Annals*, vol. 65, no. 2, pp. 621–641, 2016.

[13] R. Y. Zhong, X. Xu, E. Klotz, and S. T. Newman, "Intelligent manufacturing in the context of industry 4.0: A review," *Engineering*, vol. 3, no. 5, pp. 616–630, 2017.

[14] D. M. Tilbury, "Cyber-physical manufacturing systems," *Annual Review of Control, Robotics, and Autonomous Systems*, vol. 2, no. 1, pp. 427–443, 2019.

[15] S. Cavalieri and S. Mulè, "Interoperability between OPC UA and onem2m," *Journal of Internet Services and Applications*, vol. 12, no. 1, 2021.

[16] M. Büchter and S. Wolf, *Plug and Work with OPC UA at the Field Level: Integration of Low-Level Devices*. Springer Berlin Heidelberg, 2022, pp. 63–75.

[17] A. Fuller, Z. Fan, C. Day, and C. Barlow, "Digital twin: Enabling technologies, challenges and open research," *IEEE Access*, vol. 8, pp. 108 952–108 971, 2020.

[18] Y. Lu, C. Liu, K. I.-K. Wang, H. Huang, and X. Xu, "Digital twin-driven smart manufacturing: Connotation, reference model, applications and research issues," *Robotics and Computer-Integrated Manufacturing*, vol. 61, p. 101837, 2020.

[19] T. Schick, J. Dwivedi-Yu, R. Dessi, R. Raileanu, M. Lomeli, L. Zettlemoyer, N. Cancedda, and T. Scialom, "Toolformer: Language models can teach themselves to use tools," *arXiv preprint arXiv:2302.04761*, 2023. [Online]. Available: https://arxiv.org/abs/2302.04761

[20] S. G. Patil, T. Zhang, X. Wang, and J. E. Gonzalez, "Gorilla: Large language model connected with massive apis," *arXiv preprint arXiv:2305.15334*, 2023. [Online]. Available: https://arxiv.org/abs/2305.15334

[21] S. Qiao, H. Gui, C. Lv, Q. Jia, H. Chen, and N. Zhang, "Making language models better tool learners with execution feedback," in *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, 2024, pp. 3550–3568.